

12 information states (distinct information_state_string over both players' decision nodes)

```
=====
(1) SCENARIO   game=kuhn_poker   infoset=[1b]   player=1   legal=[0, 1]
=====
```

(2) POLICY INSPECTOR

```
student action probs   [0.676, 0.324]
CFR teacher target     [0.659, 0.341]
L1 gap vs teacher      0.0351
```

per-action logit decomposition (sums to the action score):

action	reward	opponent	intrinsic	total
0	0.375	-0.177	-0.060	0.138
1	-0.325	-0.207	-0.066	-0.598

do-intervention (zero a factor, recompute softmax):

```
remove reward    -> [0.509, 0.491]   L1 shift 0.3342
remove opponent  -> [0.670, 0.330]   L1 shift 0.0133
remove intrinsic -> [0.675, 0.325]   L1 shift 0.0028
```

(3) EVIDENCE

```
exact best-response exploitability (NashConv)   0.00779
CFR teacher exploitability                      0.00303
```

semantic diagnostics (per-head proxy check):

```
reward      vs hand strength      r=-0.084  [UNCERTAIN]
opponent    vs betting-history length r=+0.497  [VALIDATED]
intrinsic   vs state-independent by design r=  n/a  [N/A]
```

audit verdict: the reward-designated head dominates this decision

```
(proxy check FAILED -- report which head carries the logit mass,
NOT that reward structure caused the action)
```